

```
In [1]: # This Python 3 environment comes with many helpful analytics libraries installed
# It is defined by the kaggle/python Docker image: https://github.com/kaggle/docker-pyt
# For example, here's several helpful packages to load

import numpy as np # linear algebra
import pandas as pd # data processing, CSV file I/O (e.g. pd.read_csv)

# Input data files are available in the read-only "../input/" directory
# For example, running this (by clicking run or pressing Shift+Enter) will list all fil

import os
train=pd.read_csv('training.csv')
# You can write up to 20GB to the current directory (/kaggle/working/) that gets preser
# You can also write temporary files to /kaggle/temp/, but they won't be saved outside
```

Displaying the 5 rows of the dataset

```
In [2]: train.head()
```

```
Out[2]:
```

	id	name	document_text	cat_name
0	22474	Information Regarding the Merger of Navios Mar...	At a special meeting held on March 24, 2021 sh...	Corporate Communications
1	27460	Announcement on Approving the Change of Member...	On April 2, 2021, the China Financial Futures ...	Securities Settlement
2	6926	SFC Suspends Shiu Yau Wah for Five Months	The Securities and Futures Commission (SFC) ha...	Antitrust
3	6982	Renminbi RMB Haircut - February 4, 2020	Pursuant to Section 2.6.2 of the Clearing Hous...	Securities Settlement
4	5022	Anti-Money Laundering, Countering Financing of...	Money laundering and terrorism financing (ML/T...	Financial Crime

Displaying unique size of train attributes

```
In [3]: len(np.unique(list(train['document_text'])))
```

```
Out[3]: 9151
```

```
In [4]: len(np.unique(list(train['id'])))
```

```
Out[4]: 9859
```

```
In [5]: len(np.unique(list(train['name'])))
```

```
Out[5]: 8594
```

```
In [6]: len(np.unique(list((train['cat_name']))))
```

Out[6]: 50

```
In [7]: newresult = train.groupby(['id', 'name'])['document_text'].agg(list).to_dict()
# for k,v in newresult:
#     print(k,v)
```

Cat_name grouped into one array

```
In [8]: train2=train.groupby('id', sort=False).agg(lambda x: list(set(x))).reset_index()
train2.head()

train2 = train2.explode('name')

train2 = train2.explode('document_text')

train2.head()
```

```
Out[8]:
```

	id	name	document_text	cat_name
0	22474	Information Regarding the Merger of Navios Mar...	At a special meeting held on March 24, 2021 sh...	[Corporate Communications]
1	27460	Announcement on Approving the Change of Member...	On April 2, 2021, the China Financial Futures ...	[Trade Settlement, Securities Settlement]
2	6926	SFC Suspends Shiu Yau Wah for Five Months	The Securities and Futures Commission (SFC) ha...	[Compliance Management, Licensure and certific...
3	6982	Renminbi RMB Haircut - February 4, 2020	Pursuant to Section 2.6.2 of the Clearing Hous...	[Payments and Settlements, Securities Settleme...
4	5022	Anti-Money Laundering, Countering Financing of...	Money laundering and terrorism financing (ML/T...	[Regulatory Reporting, Money- Laundering and Te...

```
In [9]: # for k,v in newresult.items():
#     print(k,v)
```

```
In [10]: from sklearn.feature_extraction.text import TfidfVectorizer

# initialize the TfidfVectorizer without any parameters
tfidf_vect = TfidfVectorizer()

# with stop words removed
tfidf_vect = TfidfVectorizer(stop_words="english")

# generate tfidf matrix
dtm= tfidf_vect.fit_transform(train["document_text"])

print("type of dtm:", type(dtm))
print("size of tfidf matrix:", dtm.shape)
```

```
type of dtm: <class 'scipy.sparse.csr.csr_matrix'>
size of tfidf matrix: (47102, 33459)
```

```
In [11]: # 1. Check vocabulary

# Vocabulary is a dictionary mapping a word to an index

# the number of words in the vocabulary
print("total number of words:", len(tfidf_vect.vocabulary_))

print("type of vocabulary:", \
      type(tfidf_vect.vocabulary_))
```

total number of words: 33459
 type of vocabulary: <class 'dict'>

Displaying some of the text and most freq words

```
In [12]: # 3.4 check words with top tf-idf wights in a document,
# e.g. 1st document

# get mapping from word index to word
# i.e. reversal mapping of tfidf_vect.vocabulary_
voc_lookup={tfidf_vect.vocabulary_[word]:word \
            for word in tfidf_vect.vocabulary_}

print("\nOriginal text: \n"+train["document_text"][0])

print("\ntfidf weights: \n")

# first, covert the sparse matrix row to a dense array
doc0=dtm[0].toarray()[0]
print("Vectorized document shape: ", doc0.shape, "\n")

# get index of top 20 words
print("top words:")
top_words=(doc0.argsort()[::-1][0:20])
for i in top_words:
    print("{0}:\t{1:.3f}".format(voc_lookup[i], doc0[i]))
#[voc_lookup[i], '%.3f'%doc0[i]) for i in top_words]
```

Original text:

At a special meeting held on March 24, 2021 shareholders of Navios Maritime Containers L.P. (NMCI) approved the proposed merger with Navios Maritime Partners L.P. The merger is anticipated to become effective on March 31, 2021. The details are as follows: Company Issue: Navios Maritime Containers L.P. Common Units CUSIP#: Y62151108 Symbol: NMCI Anticipated Last Trading Date: March 31, 2021 Anticipated Marketplace Effective Date for Suspension: April 1, 2021 Merger Consideration: 0.39 of a common unit representing limited partner interests in Navios Partners for each share held.

tfidf weights:

Vectorized document shape: (33459,)

top words:

navios:	0.593
maritime:	0.389
nmci:	0.296
containers:	0.271
anticipated:	0.240
merger:	0.200
y62151108:	0.165
partners:	0.160

```

2021: 0.155
march: 0.140
held: 0.110
common: 0.108
31: 0.095
partner: 0.093
39: 0.083
representing: 0.081
date: 0.081
effective: 0.080
units: 0.072
cusip: 0.072

```

Using Multilabel binarizer to convert array of names into array of class variables

In [13]:

```

from sklearn.preprocessing import MultiLabelBinarizer
import numpy as np

```

```

mlb = MultiLabelBinarizer()
# labels=List(np.unique(train2['cat_name']))
# labels=[list(labels)]
# labels
# mlb.fit(labels)
Y=mlb.fit_transform(train2['cat_name'])
# check size of indicator matrix
# print some rows
print(Y[0:5])
print(Y.shape)
# check classes
print(mlb.classes_)

# check # of samples in each class
np.sum(Y, axis=0)

```

```

[[0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
  0 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
  0 0 0 0 0 0 0 0 0 0 0 0 1 0 1]
 [0 1 0 0 0 1 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 1 1 0 0 0 0
  0 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0
  0 0 0 0 0 0 0 1 0 0 0 1 0 1]
 [0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0
  0 0 1 1 0 0 1 0 0 0 0 0 0 0]]

```

(9859, 50)

```

['Accounting and Finance' 'Antitrust' 'Banking' 'Broker Dealer'
 'Commodities Trading' 'Compliance Management' 'Consumer protection'
 'Contract Provisions' 'Corporate Communications' 'Corporate Governance'
 'Definitions' 'Delivery' 'Examinations' 'Exemptions' 'Fees and Charges'
 'Financial Accounting' 'Financial Crime' 'Forms' 'Fraud' 'IT Risk'
 'Information Filing' 'Insurance' 'Legal' 'Legal Proceedings' 'Licensing'
 'Licensure and certification' 'Liquidity Risk' 'Listing' 'Market Abuse'
 'Market Risk' 'Monetary and Economic Policy' 'Money Services'
 'Money-Laundering and Terrorist Financing' 'Natural Disasters'
 'Payments and Settlements' 'Powers and Duties' 'Quotation'
 'Records Maintenance' 'Regulatory Actions' 'Regulatory Reporting'
 'Required Disclosures' 'Research' 'Risk Management' 'Securities Clearing'
 'Securities Issuing' 'Securities Management' 'Securities Sales'
 'Securities Settlement' 'Trade Pricing' 'Trade Settlement']

```

```

Out[13]: array([[ 935,  880, 1078,  670,  682, 1391,  969, 1153,  518,  958,  570,
    821, 1742, 1190, 1301,  535, 1178,  508,  906,  435, 1387,  737,
    907, 1343,  999,  982,  534, 1124,  722, 1633,  802,  869,  505,

```

1079, 1099, 797, 611, 630, 1621, 1042, 627, 554, 982, 1141,
1107, 664, 1737, 852, 872, 723])

In [28]:

```
import numpy as np
from sklearn.pipeline import Pipeline
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.svm import LinearSVC
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.model_selection import train_test_split
from sklearn.multiclass import OneVsRestClassifier
from sklearn.metrics import classification_report

# split dataset into train (70%) and test sets (30%)
X_train, X_test, Y_train, Y_test = train_test_split(\
    train2['document_text'], Y, test_size=0.30, random_state=0)

classifier = Pipeline([
    ('tfidf', TfidfVectorizer(stop_words="english",\
                             min_df=5)),
    ('clf', OneVsRestClassifier(LinearSVC(C=3.0)))]

classifier.fit(X_train, Y_train)
```

```
Out[28]: Pipeline(steps=[('tfidf', TfidfVectorizer(min_df=5, stop_words='english')),
                           ('clf', OneVsRestClassifier(estimator=LinearSVC(C=3.0)))])
```

Classification using SVM

In [29]:

```
from sklearn.metrics import classification_report

predicted = classifier.predict(X_test)

print(predicted.shape)
print("predicted:")
print(predicted[0:2])
print("actual:")
print(Y_test[0:2])

print(classification_report\
      (Y_test, predicted, target_names=mlb.classes_))
```

(2958, 50)

predicted:

```
[[1 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0
  0 0 0 0 0 0 0 0 0 0 0 0 0 0 0]]
[[0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0
  0 0 0 0 0 0 0 0 0 0 0 0 0 0]]
```

actual:

```
[[1 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0
  0 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 1 0 0 0 0 0 0 0
  0 0 0 0 0 0 0 0 0 0 0 0 0 0]]
```

```
precision    recall  f1-score   support
```

Accounting and Finance	0.86	0.57	0.69	294
Antitrust	0.87	0.70	0.78	265
Banking	0.93	0.77	0.84	344
Broker Dealer	0.95	0.71	0.81	220
Commodities Trading	0.93	0.75	0.83	195
Compliance Management	0.88	0.74	0.80	420

Consumer protection	0.87	0.73	0.79	286
Contract Provisions	0.87	0.71	0.78	326
Corporate Communications	0.76	0.53	0.62	177
Corporate Governance	0.73	0.52	0.61	297
Definitions	0.82	0.57	0.68	176
Delivery	0.97	0.73	0.83	241
Examinations	0.79	0.70	0.74	518
Exemptions	0.83	0.62	0.71	359
Fees and Charges	0.96	0.80	0.87	404
Financial Accounting	0.89	0.58	0.70	154
Financial Crime	0.88	0.71	0.79	349
Forms	0.90	0.65	0.75	156
Fraud	0.97	0.74	0.84	278
IT Risk	0.84	0.55	0.66	139
Information Filing	0.77	0.70	0.73	412
Insurance	0.88	0.76	0.82	228
Legal	0.68	0.51	0.58	275
Legal Proceedings	0.82	0.71	0.76	414
Licensing	0.86	0.63	0.73	328
Licensure and certification	0.83	0.64	0.73	307
Liquidity Risk	0.83	0.68	0.75	154
Listing	0.92	0.80	0.86	339
Market Abuse	0.94	0.79	0.86	220
Market Risk	0.74	0.69	0.72	465
Monetary and Economic Policy	0.82	0.64	0.72	251
Money Services	0.77	0.63	0.69	268
Money-Laundering and Terrorist Financing	0.98	0.71	0.82	149
Natural Disasters	0.94	0.80	0.86	310
Payments and Settlements	0.89	0.79	0.84	327
Powers and Duties	0.70	0.47	0.56	242
Quotation	0.90	0.72	0.80	190
Records Maintenance	0.83	0.58	0.68	187
Regulatory Actions	0.84	0.73	0.78	492
Regulatory Reporting	0.73	0.56	0.64	303
Required Disclosures	0.80	0.45	0.58	184
Research	0.81	0.64	0.71	135
Risk Management	0.82	0.69	0.75	294
Securities Clearing	0.96	0.89	0.92	352
Securities Issuing	0.85	0.71	0.77	345
Securities Management	0.93	0.73	0.82	209
Securities Sales	0.81	0.74	0.77	531
Securities Settlement	0.88	0.79	0.83	252
Trade Pricing	0.88	0.78	0.83	259
Trade Settlement	0.89	0.72	0.80	214
micro avg	0.85	0.69	0.76	14234
macro avg	0.86	0.68	0.76	14234
weighted avg	0.85	0.69	0.76	14234
samples avg	0.78	0.67	0.70	14234

/opt/conda/lib/python3.7/site-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in samples with no predicted labels. Use `zero_division` parameter to control this behavior.
_warn_prf(average, modifier, msg_start, len(result))

In [30]:

```
import numpy as np
from sklearn.pipeline import Pipeline
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.svm import LinearSVC
from sklearn.feature_extraction.text import TfidfTransformer
from sklearn.model_selection import train_test_split
from sklearn.multiclass import OneVsRestClassifier
from sklearn.metrics import classification_report
```

```

from sklearn.naive_bayes import MultinomialNB

# split dataset into train (70%) and test sets (10%)
X_train, X_test, Y_train, Y_test = train_test_split(\
    train2['document_text'], Y, test_size=0.30, random_state=0)

classifier = Pipeline([
    ('tfidf', TfidfVectorizer(stop_words="english",\
                             min_df=5)),
    ('clf', OneVsRestClassifier(MultinomialNB()))])

classifier.fit(X_train, Y_train)

```

```

Out[30]: Pipeline(steps=[('tfidf', TfidfVectorizer(min_df=5, stop_words='english')),
                          ('clf', OneVsRestClassifier(estimator=MultinomialNB()))])

```

Classification using MultinomialNB

```

In [31]: from sklearn.metrics import classification_report

predicted = classifier.predict(X_test)

print(predicted.shape)
print("predicted:")
print(predicted[0:2])
print("actual:")
print(Y_test[0:2])

print(classification_report\
      (Y_test, predicted, target_names=mlb.classes_))

```

(2958, 50)

predicted:

```

[[0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
  0 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
  0 0 0 0 0 0 0 0 0 0 0 0 0 0]]

```

actual:

```

[[1 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0 0
  0 0 0 0 0 0 0 0 0 0 0 0 0 0]
 [0 0 0 0 0 1 0 0 0 0 0 0 0 0 0 0 0 0 0 0 1 0 0 0 0 1 0 0 0 0 0 0
  0 0 0 0 0 0 0 0 0 0 0 0 0 0]]

```

	precision	recall	f1-score	support
Accounting and Finance	0.86	0.02	0.04	294
Antitrust	0.87	0.37	0.52	265
Banking	0.91	0.18	0.30	344
Broker Dealer	0.92	0.05	0.10	220
Commodities Trading	0.96	0.34	0.51	195
Compliance Management	0.97	0.21	0.35	420
Consumer protection	1.00	0.09	0.16	286
Contract Provisions	0.83	0.44	0.57	326
Corporate Communications	0.91	0.12	0.21	177
Corporate Governance	1.00	0.02	0.04	297
Definitions	1.00	0.02	0.04	176
Delivery	0.94	0.47	0.63	241
Examinations	0.84	0.32	0.47	518
Exemptions	0.98	0.11	0.20	359
Fees and Charges	0.97	0.32	0.49	404
Financial Accounting	1.00	0.05	0.09	154
Financial Crime	0.87	0.35	0.50	349
Forms	1.00	0.21	0.35	156

Fraud	0.87	0.37	0.52	278
IT Risk	1.00	0.06	0.11	139
Information Filing	0.89	0.16	0.27	412
Insurance	0.83	0.26	0.40	228
Legal	0.50	0.01	0.01	275
Legal Proceedings	0.83	0.46	0.59	414
Licensing	0.97	0.18	0.31	328
Licensure and certification	0.82	0.19	0.31	307
Liquidity Risk	0.90	0.06	0.11	154
Listing	0.92	0.30	0.46	339
Market Abuse	1.00	0.27	0.43	220
Market Risk	0.84	0.29	0.43	465
Monetary and Economic Policy	0.86	0.19	0.31	251
Money Services	0.89	0.09	0.17	268
Money-Laundering and Terrorist Financing	1.00	0.15	0.27	149
Natural Disasters	0.76	0.35	0.48	310
Payments and Settlements	0.99	0.44	0.61	327
Powers and Duties	1.00	0.00	0.01	242
Quotation	1.00	0.18	0.30	190
Records Maintenance	1.00	0.09	0.17	187
Regulatory Actions	0.79	0.50	0.61	492
Regulatory Reporting	0.88	0.07	0.13	303
Required Disclosures	0.00	0.00	0.00	184
Research	0.57	0.03	0.06	135
Risk Management	0.94	0.10	0.18	294
Securities Clearing	0.99	0.54	0.70	352
Securities Issuing	0.98	0.19	0.31	345
Securities Management	1.00	0.51	0.68	209
Securities Sales	0.86	0.40	0.54	531
Securities Settlement	0.91	0.35	0.51	252
Trade Pricing	0.76	0.27	0.40	259
Trade Settlement	0.92	0.36	0.52	214
micro avg	0.89	0.25	0.39	14234
macro avg	0.89	0.22	0.33	14234
weighted avg	0.89	0.25	0.36	14234
samples avg	0.40	0.22	0.26	14234

/opt/conda/lib/python3.7/site-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

/opt/conda/lib/python3.7/site-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in samples with no predicted labels. Use `zero_division` parameter to control this behavior.

_warn_prf(average, modifier, msg_start, len(result))

In []:

```
# Exercise 3.3.1 Grid search

# import pipeline class
from sklearn.pipeline import Pipeline
from sklearn.naive_bayes import MultinomialNB
# import GridSearch
from sklearn.model_selection import GridSearchCV

# build a pipeline which does two steps all together:
# (1) generate tfidf, and (2) train classifier
# each step is named, i.e. "tfidf", "clf"
C=[0.5,1,5,10,50,100,1000]
for i in C:
    print('C=',i)
    text_clf = Pipeline([('tfidf', TfidfVectorizer()),
                        ('clf', OneVsRestClassifier(LinearSVC(C=i)))])
```



```

    ])

# set the range of parameters to be tuned
# each parameter is defined as
# <step name>__<parameter name in step>
# e.g. min_df is a parameter of TfidfVectorizer()
# "tfidf" is the name for TfidfVectorizer()
# therefore, 'tfidf__min_df' is the parameter in grid search

parameters = {'tfidf__min_df':[1,2,4,3,5,10,6,7,8,9,15,20],
              'tfidf__stop_words':[None,"english"]}
#              'clf__C': [0.5,1.0,2.0,3,4,5,6,7,8,9,10],
#              }

# the metric used to select the best parameters
metric = "f1_macro"

# GridSearch also uses cross validation
gs_clf = GridSearchCV\
(text_clf, param_grid=parameters, \
 scoring=metric, cv=10)

# due to data volume and large parameter combinations
# it may take long time to search for optimal parameter combination
# you can use a subset of data to test
gs_clf = gs_clf.fit(X_train, Y_train)
for param_name in gs_clf.best_params_:
    print("{0}:\t{1}".format(param_name,\
                             gs_clf.best_params_[param_name]))

print("best f1 score: {:.3f}".format(gs_clf.best_score_))

```

C= 0.5

tfidf__min_df: 1

tfidf__stop_words: english

best f1 score: 0.710

C= 1

In [57]:

```

for param_name in gs_clf.best_params_:
    print("{0}:\t{1}".format(param_name,\
                             gs_clf.best_params_[param_name]))

print("best f1 score: {:.3f}".format(gs_clf.best_score_))

```

tfidf__min_df: 20

tfidf__stop_words: english

best f1 score: 0.450

In []:

Installing essential packages

In [1]:

```
pip install transformers
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Collecting transformers
  Downloading transformers-4.25.1-py3-none-any.whl (5.8 MB)
    |████████████████████████████████████████| 5.8 MB 6.8 MB/s
Requirement already satisfied: regex!=2019.12.17 in /usr/local/lib/python3.8/dist-packages (from transformers) (2022.6.2)
Requirement already satisfied: numpy>=1.17 in /usr/local/lib/python3.8/dist-packages (from transformers) (1.21.6)
Requirement already satisfied: tqdm>=4.27 in /usr/local/lib/python3.8/dist-packages (from transformers) (4.64.1)
Requirement already satisfied: requests in /usr/local/lib/python3.8/dist-packages (from transformers) (2.23.0)
Requirement already satisfied: packaging>=20.0 in /usr/local/lib/python3.8/dist-packages (from transformers) (21.3)
Requirement already satisfied: filelock in /usr/local/lib/python3.8/dist-packages (from transformers) (3.8.2)
Collecting tokenizers!=0.11.3,<0.14,>=0.11.1
  Downloading tokenizers-0.13.2-cp38-cp38-manylinux_2_17_x86_64.manylinux2014_x86_64.whl (7.6 MB)
    |████████████████████████████████████████| 7.6 MB 77.7 MB/s
Requirement already satisfied: pyyaml>=5.1 in /usr/local/lib/python3.8/dist-packages (from transformers) (6.0)
Collecting huggingface-hub<1.0,>=0.10.0
  Downloading huggingface_hub-0.11.1-py3-none-any.whl (182 kB)
    |████████████████████████████████████████| 182 kB 91.4 MB/s
Requirement already satisfied: typing-extensions>=3.7.4.3 in /usr/local/lib/python3.8/dist-packages (from huggingface-hub<1.0,>=0.10.0->transformers) (4.4.0)
Requirement already satisfied: pyparsing!=3.0.5,>=2.0.2 in /usr/local/lib/python3.8/dist-packages (from packaging>=20.0->transformers) (3.0.9)
Requirement already satisfied: urllib3!=1.25.0,!<1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.8/dist-packages (from requests->transformers) (1.24.3)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.8/dist-packages (from requests->transformers) (2.10)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.8/dist-packages (from requests->transformers) (2022.12.7)
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.8/dist-packages (from requests->transformers) (3.0.4)
Installing collected packages: tokenizers, huggingface-hub, transformers
Successfully installed huggingface-hub-0.11.1 tokenizers-0.13.2 transformers-4.25.1
```

In [2]:

```
pip install torchmetrics
```

```
Looking in indexes: https://pypi.org/simple, https://us-python.pkg.dev/colab-wheels/public/simple/
Collecting torchmetrics
  Downloading torchmetrics-0.11.0-py3-none-any.whl (512 kB)
    |████████████████████████████████████████| 512 kB 8.5 MB/s
Requirement already satisfied: numpy>=1.17.2 in /usr/local/lib/python3.8/dist-packages (from torchmetrics) (1.21.6)
Requirement already satisfied: packaging in /usr/local/lib/python3.8/dist-packages (from torchmetrics) (21.3)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.8/dist-packages (from torchmetrics) (4.4.0)
Requirement already satisfied: torch>=1.8.1 in /usr/local/lib/python3.8/dist-packages (from torchmetrics) (1.13.0+cu116)
Requirement already satisfied: pyparsing!=3.0.5,>=2.0.2 in /usr/local/lib/python3.8/dist-packages (from packaging->torchmetrics) (3.0.9)
```

Installing collected packages: torchmetrics
Successfully installed torchmetrics-0.11.0

In [3]: `pip install torchvision`

Looking in indexes: <https://pypi.org/simple>, <https://us-python.pkg.dev/colab-wheels/public/simple/>
Requirement already satisfied: torchvision in /usr/local/lib/python3.8/dist-packages (0.14.0+cu116)
Requirement already satisfied: pillow!=8.3.*,>=5.3.0 in /usr/local/lib/python3.8/dist-packages (from torchvision) (7.1.2)
Requirement already satisfied: typing-extensions in /usr/local/lib/python3.8/dist-packages (from torchvision) (4.4.0)
Requirement already satisfied: torch==1.13.0 in /usr/local/lib/python3.8/dist-packages (from torchvision) (1.13.0+cu116)
Requirement already satisfied: numpy in /usr/local/lib/python3.8/dist-packages (from torchvision) (1.21.6)
Requirement already satisfied: requests in /usr/local/lib/python3.8/dist-packages (from torchvision) (2.23.0)
Requirement already satisfied: idna<3,>=2.5 in /usr/local/lib/python3.8/dist-packages (from requests->torchvision) (2.10)
Requirement already satisfied: urllib3!=1.25.0,!1.25.1,<1.26,>=1.21.1 in /usr/local/lib/python3.8/dist-packages (from requests->torchvision) (1.24.3)
Requirement already satisfied: certifi>=2017.4.17 in /usr/local/lib/python3.8/dist-packages (from requests->torchvision) (2022.12.7)
Requirement already satisfied: chardet<4,>=3.0.2 in /usr/local/lib/python3.8/dist-packages (from requests->torchvision) (3.0.4)

In [4]: `import random, pickle
import numpy as np
from torch.nn import BCEWithLogitsLoss, BCELoss
from sklearn.metrics import classification_report, confusion_matrix, multilabel_confusion_matrix
import tensorflow as tf
import torch
import pandas as pd
from torchmetrics.classification import MultilabelF1Score

from transformers import AutoConfig, AutoModel, AutoTokenizer, AutoModelForSequenceClassification

import torch.nn as nn
import torch.nn.functional as F
from torch.utils.data import TensorDataset, random_split
from torch.utils.data import DataLoader, RandomSampler, SequentialSampler

import copy
from sklearn.utils import shuffle
import glob

import time
import datetime

import pandas as pd
import numpy as np
import io`

In [5]: `# If there's a GPU available...
if torch.cuda.is_available():

 # Tell PyTorch to use the GPU.`

```

device = torch.device("cuda")

print('There are %d GPU(s) available.' % torch.cuda.device_count())

print('We will use the GPU:', torch.cuda.get_device_name(0))

# If not...
else:
    print('No GPU available, using the CPU instead.')
    device = torch.device("cpu")

```

There are 1 GPU(s) available.
We will use the GPU: A100-SXM4-40GB

```

In [6]: from google.colab import drive

drive.mount('/content/gdrive/', force_remount=True)

```

Mounted at /content/gdrive/

```

In [7]: path = "/content/gdrive/MyDrive/BIA667_FinalProject/train.csv"

```

```

In [8]: # from google.colab import files

# uploaded = files.upload()

```

```

In [9]: from sklearn.preprocessing import MultiLabelBinarizer
from sklearn.pipeline import Pipeline
from sklearn.feature_extraction.text import CountVectorizer

from sklearn.model_selection import train_test_split

```

```

In [10]: train_df = pd.read_csv(path)
# train_df = pd.read_csv('C:/Users/Admin/Documents/Stevens Docs/BIA 656/Project/train.c

```

Preprocessing data i.e is removing spec characters from the ttext so we have only words in the document

```

In [13]: train_df['document_text'] = train_df['document_text'].str.lower()
train_df['name'] = train_df['name'].str.lower()

```

```

In [14]: spec_chars = ["!", "'", "#", "%", "&", "(", ")", " ",
                        "*", "+", ",", "-", ".", "/", ":", ";", "<",
                        "=", ">", "?", "@", "[", "\\", "]", "^", "_",
                        "`", "{", "|", "}", "~", "-"]

for char in spec_chars:
    train_df['document_text'] = train_df['document_text'].str.replace(char, ' ')
    train_df['name'] = train_df['name'].str.replace(char, ' ')

```

<ipython-input-14-4b1c388cc09c>:6: FutureWarning: The default value of regex will change from True to False in a future version. In addition, single character regular expression

```
s will *not* be treated as literal strings when regex=True.
train_df['document_text'] = train_df['document_text'].str.replace(char, ' ')
<ipython-input-14-4b1c388cc09c>:7: FutureWarning: The default value of regex will change
from True to False in a future version. In addition, single character regular expression
s will *not* be treated as literal strings when regex=True.
train_df['name'] = train_df['name'].str.replace(char, ' ')
```

```
In [15]: df2 = train_df.groupby(['name', 'document_text'])['cat_name'].apply(list)
df3=df2.to_frame().reset_index()
```

```
In [16]: df3['comb_name_doctext'] = df3['name'] + df3['document_text'].apply(lambda x: ' ' + x)
```

```
In [17]: df3.head()
```

```
Out[17]:
```

	name	document_text	cat_name	comb_name_doctext
0	correction to symbol information regarding t...	qutoutiao inc qtt will effect a one for ten...	[Broker Dealer]	correction to symbol information regarding t...
1	digital health 2020 eu on the move wojci...	european data protection supervisor published ...	[Research, Natural Disasters, Powers and Dutie...	digital health 2020 eu on the move wojci...
2	updated correction to merger consideration ...	the business combination of quidel corp qdel ...	[Broker Dealer, Corporate Communications]	updated correction to merger consideration ...
3	updated information regarding the business c...	the business combination of twc tech holdings ...	[Forms, Listing, Broker Dealer, Securities Set...	updated information regarding the business c...
4	updated closed information regarding the bus...	the business combination of mountain crest acq...	[Broker Dealer, Corporate Communications]	updated closed information regarding the bus...

```
In [18]: len(df3)
```

```
Out[18]: 9263
```

Converting cat_name to array of binary class variables

```
In [19]: mlb = MultiLabelBinarizer()
labels = df3['cat_name'].values
label_onehot = mlb.fit_transform(labels)
```

BERT for pretrained word vectors

```
In [20]: # Load the BERT tokenizer.
print('Loading BERT tokenizer...')
#tokenizer = BertTokenizer.from_pretrained('bert-base-uncased', do_lower_case=True)

tokenizer = AutoTokenizer.from_pretrained(
```

```
"nlpueb/legal-bert-base-uncased"
```

```
)
```

Loading BERT tokenizer...

In [21]:

```
bert_model = AutoModel.from_pretrained(
    # "bert-base-uncased",
    "nlpueb/legal-bert-base-uncased",
    output_attentions = False, # Whether the model returns attentions weights.
    output_hidden_states = True, # Whether the model returns all hidden-states.

)
bert_model.cuda()
```

Some weights of the model checkpoint at nlpueb/legal-bert-base-uncased were not used when initializing BertModel: ['cls.predictions.transform.dense.weight', 'cls.predictions.transform.LayerNorm.bias', 'cls.predictions.decoder.bias', 'cls.predictions.bias', 'cls.seq_relationship.weight', 'cls.seq_relationship.bias', 'cls.predictions.transform.dense.bias', 'cls.predictions.transform.LayerNorm.weight', 'cls.predictions.decoder.weight']

- This IS expected if you are initializing BertModel from the checkpoint of a model trained on another task or with another architecture (e.g. initializing a BertForSequenceClassification model from a BertForPreTraining model).
- This IS NOT expected if you are initializing BertModel from the checkpoint of a model that you expect to be exactly identical (initializing a BertForSequenceClassification model from a BertForSequenceClassification model).

Out[21]:

```
BertModel(
  (embeddings): BertEmbeddings(
    (word_embeddings): Embedding(30522, 768, padding_idx=0)
    (position_embeddings): Embedding(512, 768)
    (token_type_embeddings): Embedding(2, 768)
    (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
    (dropout): Dropout(p=0.1, inplace=False)
  )
  (encoder): BertEncoder(
    (layer): ModuleList(
      (0): BertLayer(
        (attention): BertAttention(
          (self): BertSelfAttention(
            (query): Linear(in_features=768, out_features=768, bias=True)
            (key): Linear(in_features=768, out_features=768, bias=True)
            (value): Linear(in_features=768, out_features=768, bias=True)
            (dropout): Dropout(p=0.1, inplace=False)
          )
          (output): BertSelfOutput(
            (dense): Linear(in_features=768, out_features=768, bias=True)
            (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
            (dropout): Dropout(p=0.1, inplace=False)
          )
        )
      )
      (intermediate): BertIntermediate(
        (dense): Linear(in_features=768, out_features=3072, bias=True)
        (intermediate_act_fn): GELUActivation()
      )
      (output): BertOutput(
        (dense): Linear(in_features=3072, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
      )
    )
  )
)
```

```

    )
  )
  (1): BertLayer(
    (attention): BertAttention(
      (self): BertSelfAttention(
        (query): Linear(in_features=768, out_features=768, bias=True)
        (key): Linear(in_features=768, out_features=768, bias=True)
        (value): Linear(in_features=768, out_features=768, bias=True)
        (dropout): Dropout(p=0.1, inplace=False)
      )
      (output): BertSelfOutput(
        (dense): Linear(in_features=768, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
      )
    )
    (intermediate): BertIntermediate(
      (dense): Linear(in_features=768, out_features=3072, bias=True)
      (intermediate_act_fn): GELUActivation()
    )
    (output): BertOutput(
      (dense): Linear(in_features=3072, out_features=768, bias=True)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
  (2): BertLayer(
    (attention): BertAttention(
      (self): BertSelfAttention(
        (query): Linear(in_features=768, out_features=768, bias=True)
        (key): Linear(in_features=768, out_features=768, bias=True)
        (value): Linear(in_features=768, out_features=768, bias=True)
        (dropout): Dropout(p=0.1, inplace=False)
      )
      (output): BertSelfOutput(
        (dense): Linear(in_features=768, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
      )
    )
    (intermediate): BertIntermediate(
      (dense): Linear(in_features=768, out_features=3072, bias=True)
      (intermediate_act_fn): GELUActivation()
    )
    (output): BertOutput(
      (dense): Linear(in_features=3072, out_features=768, bias=True)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
  (3): BertLayer(
    (attention): BertAttention(
      (self): BertSelfAttention(
        (query): Linear(in_features=768, out_features=768, bias=True)
        (key): Linear(in_features=768, out_features=768, bias=True)
        (value): Linear(in_features=768, out_features=768, bias=True)
        (dropout): Dropout(p=0.1, inplace=False)
      )
      (output): BertSelfOutput(
        (dense): Linear(in_features=768, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
      )
    )
    (intermediate): BertIntermediate(

```

```

        (dense): Linear(in_features=768, out_features=3072, bias=True)
        (intermediate_act_fn): GELUActivation()
    )
    (output): BertOutput(
        (dense): Linear(in_features=3072, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
    )
)
(4): BertLayer(
    (attention): BertAttention(
        (self): BertSelfAttention(
            (query): Linear(in_features=768, out_features=768, bias=True)
            (key): Linear(in_features=768, out_features=768, bias=True)
            (value): Linear(in_features=768, out_features=768, bias=True)
            (dropout): Dropout(p=0.1, inplace=False)
        )
        (output): BertSelfOutput(
            (dense): Linear(in_features=768, out_features=768, bias=True)
            (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
            (dropout): Dropout(p=0.1, inplace=False)
        )
    )
    (intermediate): BertIntermediate(
        (dense): Linear(in_features=768, out_features=3072, bias=True)
        (intermediate_act_fn): GELUActivation()
    )
    (output): BertOutput(
        (dense): Linear(in_features=3072, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
    )
)
(5): BertLayer(
    (attention): BertAttention(
        (self): BertSelfAttention(
            (query): Linear(in_features=768, out_features=768, bias=True)
            (key): Linear(in_features=768, out_features=768, bias=True)
            (value): Linear(in_features=768, out_features=768, bias=True)
            (dropout): Dropout(p=0.1, inplace=False)
        )
        (output): BertSelfOutput(
            (dense): Linear(in_features=768, out_features=768, bias=True)
            (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
            (dropout): Dropout(p=0.1, inplace=False)
        )
    )
    (intermediate): BertIntermediate(
        (dense): Linear(in_features=768, out_features=3072, bias=True)
        (intermediate_act_fn): GELUActivation()
    )
    (output): BertOutput(
        (dense): Linear(in_features=3072, out_features=768, bias=True)
        (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
        (dropout): Dropout(p=0.1, inplace=False)
    )
)
(6): BertLayer(
    (attention): BertAttention(
        (self): BertSelfAttention(
            (query): Linear(in_features=768, out_features=768, bias=True)
            (key): Linear(in_features=768, out_features=768, bias=True)
            (value): Linear(in_features=768, out_features=768, bias=True)
            (dropout): Dropout(p=0.1, inplace=False)
        )

```



```

        (output): BertSelfOutput(
          (dense): Linear(in_features=768, out_features=768, bias=True)
          (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
          (dropout): Dropout(p=0.1, inplace=False)
        )
      )
    (intermediate): BertIntermediate(
      (dense): Linear(in_features=768, out_features=3072, bias=True)
      (intermediate_act_fn): GELUActivation()
    )
    (output): BertOutput(
      (dense): Linear(in_features=3072, out_features=768, bias=True)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
(7): BertLayer(
  (attention): BertAttention(
    (self): BertSelfAttention(
      (query): Linear(in_features=768, out_features=768, bias=True)
      (key): Linear(in_features=768, out_features=768, bias=True)
      (value): Linear(in_features=768, out_features=768, bias=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (output): BertSelfOutput(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
  (intermediate): BertIntermediate(
    (dense): Linear(in_features=768, out_features=3072, bias=True)
    (intermediate_act_fn): GELUActivation()
  )
  (output): BertOutput(
    (dense): Linear(in_features=3072, out_features=768, bias=True)
    (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
    (dropout): Dropout(p=0.1, inplace=False)
  )
)
(8): BertLayer(
  (attention): BertAttention(
    (self): BertSelfAttention(
      (query): Linear(in_features=768, out_features=768, bias=True)
      (key): Linear(in_features=768, out_features=768, bias=True)
      (value): Linear(in_features=768, out_features=768, bias=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (output): BertSelfOutput(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
  (intermediate): BertIntermediate(
    (dense): Linear(in_features=768, out_features=3072, bias=True)
    (intermediate_act_fn): GELUActivation()
  )
  (output): BertOutput(
    (dense): Linear(in_features=3072, out_features=768, bias=True)
    (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
    (dropout): Dropout(p=0.1, inplace=False)
  )
)
(9): BertLayer(

```

```

(attention): BertAttention(
  (self): BertSelfAttention(
    (query): Linear(in_features=768, out_features=768, bias=True)
    (key): Linear(in_features=768, out_features=768, bias=True)
    (value): Linear(in_features=768, out_features=768, bias=True)
    (dropout): Dropout(p=0.1, inplace=False)
  )
  (output): BertSelfOutput(
    (dense): Linear(in_features=768, out_features=768, bias=True)
    (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
    (dropout): Dropout(p=0.1, inplace=False)
  )
)
(intermediate): BertIntermediate(
  (dense): Linear(in_features=768, out_features=3072, bias=True)
  (intermediate_act_fn): GELUActivation()
)
(output): BertOutput(
  (dense): Linear(in_features=3072, out_features=768, bias=True)
  (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
  (dropout): Dropout(p=0.1, inplace=False)
)
)
(10): BertLayer(
  (attention): BertAttention(
    (self): BertSelfAttention(
      (query): Linear(in_features=768, out_features=768, bias=True)
      (key): Linear(in_features=768, out_features=768, bias=True)
      (value): Linear(in_features=768, out_features=768, bias=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (output): BertSelfOutput(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
  (intermediate): BertIntermediate(
    (dense): Linear(in_features=768, out_features=3072, bias=True)
    (intermediate_act_fn): GELUActivation()
  )
  (output): BertOutput(
    (dense): Linear(in_features=3072, out_features=768, bias=True)
    (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
    (dropout): Dropout(p=0.1, inplace=False)
  )
)
)
(11): BertLayer(
  (attention): BertAttention(
    (self): BertSelfAttention(
      (query): Linear(in_features=768, out_features=768, bias=True)
      (key): Linear(in_features=768, out_features=768, bias=True)
      (value): Linear(in_features=768, out_features=768, bias=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
    (output): BertSelfOutput(
      (dense): Linear(in_features=768, out_features=768, bias=True)
      (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
      (dropout): Dropout(p=0.1, inplace=False)
    )
  )
  (intermediate): BertIntermediate(
    (dense): Linear(in_features=768, out_features=3072, bias=True)
    (intermediate_act_fn): GELUActivation()
  )
)

```

```

        (output): BertOutput(
          (dense): Linear(in_features=3072, out_features=768, bias=True)
          (LayerNorm): LayerNorm((768,), eps=1e-12, elementwise_affine=True)
          (dropout): Dropout(p=0.1, inplace=False)
        )
      )
    )
  )
  (pooler): BertPooler(
    (dense): Linear(in_features=768, out_features=768, bias=True)
    (activation): Tanh()
  )
)

```

In [23]:

```

import torch
from torch.utils.data import TensorDataset, DataLoader

def get_pretrained_word_vectors(sentences, tokenizer, bert_model):
    """
    Obtain the pretrained word vectors for the given sentences using a BERT model.

    Parameters:
        - sentences (list): List of sentences to process
        - tokenizer (object): Tokenizer object to use for encoding the sentences
        - bert_model (object): BERT model to use for generating the word vectors

    Returns:
        - token_embeddings (list): List of token embeddings for the input sentences
        - attention_masks (torch.Tensor): Tensor of attention masks for the input sentences
        - input_ids (torch.Tensor): Tensor of input IDs for the input sentences
    """
    input_ids = []
    attention_masks = []

    # Tokenize each sentence and create input IDs and attention masks
    for sent in sentences:
        # Tokenize and create special tokens, pad or truncate, and create attention masks
        encoded_dict = tokenizer.encode_plus(
            sent, # Sentence to encode
            add_special_tokens=True, # Add '[CLS]' and '[SEP]'
            max_length=400, # Pad & truncate all sentences
            truncation=True,
            padding='max_length',
            return_attention_mask=True, # Construct attention masks
            return_tensors='pt', # Return pytorch tensors
        )

        # Append the encoded sentence and attention mask to the lists
        input_ids.append(encoded_dict['input_ids'])
        attention_masks.append(encoded_dict['attention_mask'])

    # Convert lists to tensors
    input_ids = torch.cat(input_ids, dim=0)
    attention_masks = torch.cat(attention_masks, dim=0)

    # Create a TensorDataset and DataLoader for the input IDs and attention masks
    dataset = TensorDataset(input_ids, attention_masks)
    data_loader = DataLoader(dataset, batch_size=8, num_workers=2)

    token_embeddings = []

```

```

# Set the model to evaluation mode
bert_model.eval()
with torch.no_grad():
    # Iterate over the data in the data loader
    for input_id, attention_mask in data_loader:
        # Get the hidden states of the model for the input IDs and attention masks
        outputs = bert_model(input_id.to(device), attention_mask.to(device))
        hidden_states = outputs[2]

        # Stack the last four hidden states and permute the dimensions
        emb = torch.stack(hidden_states[-4:], dim=0)
        emb = emb.permute(1, 2, 0, 3)

        # Take the mean of the last four hidden states along the third dimension
        emb = emb.mean(axis=2)

        # Append the token embeddings to the list
        # Append the token embeddings to the list
        token_embeddings.append(emb)

    return token_embeddings, attention_masks, input_ids

```

```
In [24]: token_embeddings, masks, input_ids = get_pretrained_word_vectors(df3['comb_name_doctext'])
```

```
In [25]: token_embeddings_splice = np.array([j.cpu().numpy() for i in token_embeddings for j in
token_embeddings_masked = (masks.view(-1, 400, 1) * token_embeddings_splice)
```

Dividing the data into train and test set with Y labels

```
In [27]: X_train, X_test, Y_train, Y_test = train_test_split(token_embeddings_masked, label_oneh
```

```
In [28]: from torch.utils.data import TensorDataset, random_split, Dataset
from torch.utils.data import DataLoader, RandomSampler, SequentialSampler
```

Creating Custom Dataset for it to be used in Neural Network

```
In [29]: class Text_dataset(Dataset):
    def __init__(self, features, labels):
        self.length = len(labels)
        self.features = torch.Tensor(features)
        self.labels = torch.Tensor(labels)

    def __getitem__(self, index):
        return self.features[index], self.labels[index]

    def __len__(self):
        return self.length
```

```
In [30]: # datasets
train_dataset = Text_dataset(X_train, Y_train)
test_dataset = Text_dataset(X_test, Y_test)
```

```
In [31]: train_dataset.features.size()
```

```
Out[31]: torch.Size([7410, 400, 768])
```

Defining the RNN Model class which has unigram,bigram,trigram and classifier layers

```
In [32]: class TextRNN(nn.Module):
    def __init__(self, DOC_LEN, embedding_dim, dropout_ratio):
        super(TextRNN, self).__init__()
        self.dropout_ratio = dropout_ratio
        in_channels = 400
        out_channels = 600
        bias = False

        # reduce the length
        self.reduce = nn.Sequential(
            nn.Linear(in_features=embedding_dim, out_features=in_channels),
            # nn.Dropout(dropout_ratio),
            nn.ReLU()
        )

        # unigram RNN
        self.unigram = nn.GRU(input_size=in_channels, hidden_size=out_channels, num_layers=1)

        # bigram RNN
        self.bigram = nn.GRU(input_size=in_channels, hidden_size=out_channels, num_layers=1)

        # trigram RNN
        self.trigram = nn.GRU(input_size=in_channels, hidden_size=out_channels, num_layers=1)

        # simple classifier
        self.classifier = nn.Sequential(
            nn.Dropout(dropout_ratio),
            nn.Linear(in_features=out_channels*3, out_features=50)
        )

    def forward(self, x):
        # reduce length
        x = self.reduce(x)
        # print(x.shape)
        x = torch.transpose(x, dim0=1, dim1=2) # (-1, DOC_LEN, embedding_dim): embedding
        # print(x.shape)
        # unigram RNN output
        uni_gram_output, _ = self.unigram(x)
        uni_gram_output = uni_gram_output[:, -1, :]

        # bigram RNN output
        bi_gram_output, _ = self.bigram(x)
        bi_gram_output = bi_gram_output[:, -1, :]

        # trigram RNN output
        tri_gram_output, _ = self.trigram(x)
        tri_gram_output = tri_gram_output[:, -1, :]

        # concatenate
        x = torch.cat((uni_gram_output,
                        bi_gram_output,
```

```

        tri_gram_output
        ), dim=1)

    # classifier
    x = self.classifier(x)

    return x

```

```
In [33]: model_RNN=TextRNN(400,768,0.5)
```

```
In [34]: # Compute weights to handle imbalanced classes
class_counts = Y_train.sum(axis=0)
weights = [1/count for count in class_counts]
weights = [weight/sum(weights) for weight in weights]
weights = torch.Tensor(weights).to(device)
```

Defining train function for the model

```
In [35]: def train_model(model, train_dataset, test_dataset, device, lr=0.0001, epochs=20, batch
# construct dataloader
train_loader = DataLoader(train_dataset, batch_size=batch_size, shuffle=True)
test_loader = DataLoader(test_dataset, batch_size=batch_size, shuffle=True)
test_pred = DataLoader(test_dataset, batch_size=len(test_dataset))

# move model to device
model = model.to(device)

# history
history = {'train_loss': [],
          'train_f1': [],
          'test_loss': [],
          'test_f1': []
          }

# setup loss function and optimizer
criterion = nn.BCEWithLogitsLoss(reduction='none')
optimizer = torch.optim.Adam(model.parameters(), lr=lr)
metric = MultilabelF1Score(num_labels=50).to(device)

# training loop
print('Training Start')
for epoch in range(epochs):
    model.train()
    train_loss = 0
    train_f1 = 0
    test_loss = 0
    test_f1 = 0
    for x, y in train_loader:
        # move data to device
        x = x.to(device)
        y = y.to(device)
        # forward
        outputs = model(x.float())
        pred = torch.round(torch.sigmoid(outputs))
        cur_train_loss = criterion(outputs, y)
        cur_train_loss = (cur_train_loss * weights).mean()*100
        cur_train_f1 = metric(pred, y)
    # backward

```

```

        cur_train_loss.backward()
        optimizer.step()
        optimizer.zero_grad()
        # loss and acc
        train_loss += cur_train_loss
        train_f1 += cur_train_f1

    # test start
    model.eval()
    with torch.no_grad():
        for x, y in test_loader:
            # move
            x = x.to(device)
            y = y.to(device)
            # predict
            outputs = model(x.float())
            pred = torch.round(torch.sigmoid(outputs))
            cur_test_loss = criterion(outputs, y).mean()
            cur_test_loss = (cur_train_loss * weights).mean()*100
            cur_test_f1 = metric(pred, y)
            # loss and acc
            test_loss += cur_test_loss
            test_f1 += cur_test_f1

    # epoch output
    train_loss = (train_loss/len(train_loader)).item()
    train_f1 = (train_f1/len(train_loader)).item()
    val_loss = (test_loss/len(test_loader)).item()
    val_f1 = (test_f1/len(test_loader)).item()
    history['train_loss'].append(train_loss)
    history['train_f1'].append(train_f1)
    history['test_loss'].append(val_loss)
    history['test_f1'].append(val_f1)
    print(f"Epoch:{epoch + 1} / {epochs}, train loss:{train_loss:.4f} train f1:{tra

with torch.no_grad():
    for x, y in test_pred:
        x = x.to(device)
        y = y.to(device)

        outputs = model(x.float())
        pred = torch.round(torch.sigmoid(outputs))
        preds = pred.cpu().detach().numpy()

return history, preds

```

In [36]:

```

history, preds = train_model(model=model_RNN,
                             train_dataset=train_dataset,
                             test_dataset=test_dataset,
                             device=device,
                             lr=0.0005,
                             epochs=40,
                             batch_size=128)

```

Training Start

```

Epoch:1 / 40, train loss:0.6253 train f1:0.0038, valid loss:1.1084 valid f1:0.0000
Epoch:2 / 40, train loss:0.5372 train f1:0.0049, valid loss:0.9878 valid f1:0.0096
Epoch:3 / 40, train loss:0.4938 train f1:0.0622, valid loss:0.9451 valid f1:0.1249

```

```

Epoch:4 / 40, train loss:0.4658 train f1:0.1375, valid loss:0.8998 valid f1:0.1615
Epoch:5 / 40, train loss:0.4440 train f1:0.1906, valid loss:0.9001 valid f1:0.1959
Epoch:6 / 40, train loss:0.4225 train f1:0.2403, valid loss:0.8345 valid f1:0.2467
Epoch:7 / 40, train loss:0.4011 train f1:0.3019, valid loss:0.7973 valid f1:0.2943
Epoch:8 / 40, train loss:0.3804 train f1:0.3566, valid loss:0.7716 valid f1:0.3578
Epoch:9 / 40, train loss:0.3592 train f1:0.4102, valid loss:0.6918 valid f1:0.3873
Epoch:10 / 40, train loss:0.3423 train f1:0.4531, valid loss:0.7030 valid f1:0.3948
Epoch:11 / 40, train loss:0.3221 train f1:0.4983, valid loss:0.6898 valid f1:0.4171
Epoch:12 / 40, train loss:0.3038 train f1:0.5337, valid loss:0.6222 valid f1:0.4537
Epoch:13 / 40, train loss:0.2822 train f1:0.5780, valid loss:0.5735 valid f1:0.4620
Epoch:14 / 40, train loss:0.2624 train f1:0.6123, valid loss:0.5143 valid f1:0.4806
Epoch:15 / 40, train loss:0.2450 train f1:0.6471, valid loss:0.4902 valid f1:0.4968
Epoch:16 / 40, train loss:0.2229 train f1:0.6869, valid loss:0.4142 valid f1:0.4978
Epoch:17 / 40, train loss:0.2048 train f1:0.7187, valid loss:0.3707 valid f1:0.5149
Epoch:18 / 40, train loss:0.1858 train f1:0.7538, valid loss:0.3747 valid f1:0.5120
Epoch:19 / 40, train loss:0.1693 train f1:0.7797, valid loss:0.3213 valid f1:0.5088
Epoch:20 / 40, train loss:0.1511 train f1:0.8087, valid loss:0.3093 valid f1:0.5226
Epoch:21 / 40, train loss:0.1365 train f1:0.8329, valid loss:0.2539 valid f1:0.5099
Epoch:22 / 40, train loss:0.1212 train f1:0.8544, valid loss:0.2619 valid f1:0.5223
Epoch:23 / 40, train loss:0.1070 train f1:0.8755, valid loss:0.1923 valid f1:0.5237
Epoch:24 / 40, train loss:0.0943 train f1:0.8944, valid loss:0.1921 valid f1:0.5263
Epoch:25 / 40, train loss:0.0839 train f1:0.9104, valid loss:0.1546 valid f1:0.5095
Epoch:26 / 40, train loss:0.0735 train f1:0.9234, valid loss:0.1646 valid f1:0.5235
Epoch:27 / 40, train loss:0.0656 train f1:0.9347, valid loss:0.1426 valid f1:0.5173
Epoch:28 / 40, train loss:0.0592 train f1:0.9421, valid loss:0.1163 valid f1:0.5203
Epoch:29 / 40, train loss:0.0523 train f1:0.9508, valid loss:0.1046 valid f1:0.5142
Epoch:30 / 40, train loss:0.0475 train f1:0.9558, valid loss:0.0949 valid f1:0.5225
Epoch:31 / 40, train loss:0.0417 train f1:0.9638, valid loss:0.0809 valid f1:0.5210
Epoch:32 / 40, train loss:0.0382 train f1:0.9671, valid loss:0.0759 valid f1:0.5234
Epoch:33 / 40, train loss:0.0335 train f1:0.9726, valid loss:0.0631 valid f1:0.5085
Epoch:34 / 40, train loss:0.0295 train f1:0.9780, valid loss:0.0642 valid f1:0.5137
Epoch:35 / 40, train loss:0.0266 train f1:0.9803, valid loss:0.0500 valid f1:0.5147
Epoch:36 / 40, train loss:0.0239 train f1:0.9829, valid loss:0.0435 valid f1:0.5275
Epoch:37 / 40, train loss:0.0218 train f1:0.9847, valid loss:0.0489 valid f1:0.5130
Epoch:38 / 40, train loss:0.0192 train f1:0.9878, valid loss:0.0369 valid f1:0.5113
Epoch:39 / 40, train loss:0.0183 train f1:0.9884, valid loss:0.0437 valid f1:0.5234
Epoch:40 / 40, train loss:0.0167 train f1:0.9893, valid loss:0.0405 valid f1:0.5252

```

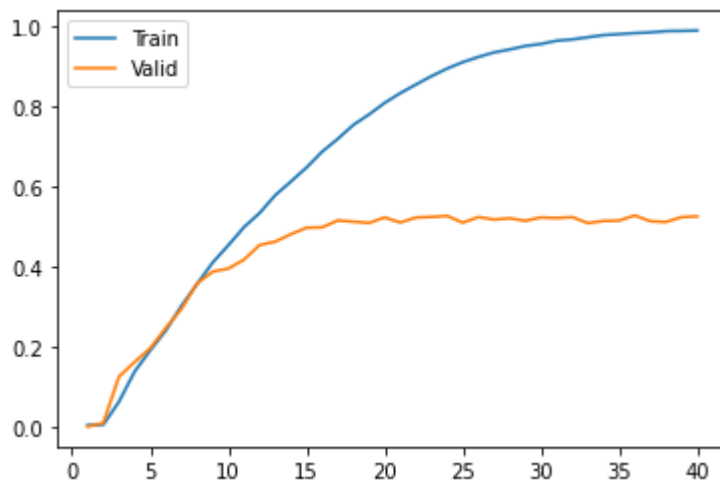
Plot for F1 score for RNN model

```

In [37]: import matplotlib.pyplot as plt
          plt.plot(range(1, 41), history['train_f1'], label='Train')
          plt.plot(range(1, 41), history['test_f1'], label='Valid')
          plt.legend()
          plt.plot()

```

Out[37]: []



In [38]:

```
#ModelLRNN
from sklearn.metrics import classification_report

print(classification_report(Y_test, preds))
```

	precision	recall	f1-score	support
0	0.60	0.41	0.49	191
1	0.79	0.64	0.71	171
2	0.55	0.49	0.52	194
3	0.67	0.48	0.56	134
4	0.74	0.67	0.70	104
5	0.67	0.46	0.54	264
6	0.69	0.39	0.49	187
7	0.78	0.58	0.67	219
8	0.70	0.29	0.41	112
9	0.56	0.26	0.36	177
10	0.34	0.20	0.25	102
11	0.81	0.62	0.70	147
12	0.66	0.46	0.54	344
13	0.53	0.37	0.43	229
14	0.55	0.51	0.53	223
15	0.62	0.43	0.51	109
16	0.74	0.60	0.66	218
17	0.76	0.32	0.45	101
18	0.71	0.65	0.68	165
19	0.43	0.29	0.34	73
20	0.56	0.43	0.49	261
21	0.71	0.56	0.62	149
22	0.45	0.17	0.24	175
23	0.71	0.47	0.57	259
24	0.55	0.40	0.46	184
25	0.59	0.38	0.46	178
26	0.63	0.39	0.48	101
27	0.74	0.55	0.63	203
28	0.87	0.65	0.75	138
29	0.70	0.53	0.60	319
30	0.59	0.35	0.44	152
31	0.51	0.38	0.43	169
32	0.81	0.54	0.65	89
33	0.68	0.61	0.64	199
34	0.85	0.59	0.70	200
35	0.32	0.08	0.13	143
36	0.67	0.58	0.62	111
37	0.59	0.34	0.43	112
38	0.73	0.53	0.61	304

39	0.61	0.41	0.49	179
40	0.46	0.24	0.32	134
41	0.44	0.18	0.26	121
42	0.61	0.55	0.58	172
43	0.88	0.76	0.82	204
44	0.66	0.47	0.55	192
45	0.82	0.62	0.70	123
46	0.70	0.66	0.68	308
47	0.77	0.60	0.67	152
48	0.68	0.54	0.60	157
49	0.84	0.60	0.70	144
micro avg	0.67	0.48	0.56	8796
macro avg	0.65	0.47	0.54	8796
weighted avg	0.66	0.48	0.55	8796
samples avg	0.58	0.45	0.48	8796

```
/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in samples with no predicted labels. Use `zero_division` parameter to control this behavior.
  _warn_prf(average, modifier, msg_start, len(result))
```

Defining advanced model with CNN with unigram,bigram,trigram,quadgram,pentagram and one convolution layer with max pooling

In [39]:

```
class TextCNN(nn.Module):
    def __init__(self, DOC_LEN, embedding_dim, dropout_ratio):
        super(TextCNN, self).__init__()
        self.DOC_LEN = DOC_LEN
        self.in_channels = 400
        self.out_channels = 600
        self.bias = False

        # reduce the length
        self.reduce_length = nn.Sequential(
            nn.Linear(in_features=embedding_dim, out_features=self.in_channels),
            nn.ReLU()
        )

        # 1D CNNs for unigrams, bigrams, trigrams, quadgrams, and pentagrams
        self.unigram_cnn = self._build_1d_cnn(kernel_size=1)
        self.bigram_cnn = self._build_1d_cnn(kernel_size=2)
        self.trigram_cnn = self._build_1d_cnn(kernel_size=3)
        self.quadgram_cnn = self._build_1d_cnn(kernel_size=4)
        self.pentagram_cnn = self._build_1d_cnn(kernel_size=5)

        # simple classifier
        self.classifier = nn.Sequential(
            nn.Dropout(dropout_ratio),
            nn.Linear(in_features=self.out_channels*5, out_features=50)
        )

    def _build_1d_cnn(self, kernel_size):
        """Helper function to build 1D CNN layers for a given kernel size."""
        return nn.Sequential(
            nn.Conv1d(in_channels=self.in_channels, out_channels=self.out_channels, kernel_size=kernel_size),
            nn.ReLU(),
            nn.MaxPool1d(kernel_size=self.DOC_LEN - kernel_size + 1),
            nn.Flatten()
        )
```

```

def forward(self, x):
    # reduce length
    x = self.reduce_length(x)
    x = torch.transpose(x, dim0=1, dim1=2) # (-1, DOC_LEN, embedding_dim)

    # 1D CNN outputs
    unigram_output = self.unigram_cnn(x)
    bigram_output = self.bigram_cnn(x)
    trigram_output = self.trigram_cnn(x)
    quadgram_output = self.quadgram_cnn(x)
    pentagram_output = self.pentagram_cnn(x)

    # concatenate 1D CNN outputs
    x = torch.cat((unigram_output, bigram_output, trigram_output, quadgram_output,
                    pentagram_output))

    # classifier
    x = self.classifier(x)

    return x

```

```
In [40]: model_CNN=TextCNN(400,768,0.5)
```

```
In [41]: history, preds = train_model(model=model_CNN,
                                     train_dataset=train_dataset,
                                     test_dataset=test_dataset,
                                     device=device,
                                     lr=0.0005,
                                     epochs=40,
                                     batch_size=128)
```

Training Start

```

Epoch:1 / 40, train loss:0.5925 train f1:0.0026, valid loss:1.0962 valid f1:0.0000
Epoch:2 / 40, train loss:0.5010 train f1:0.0495, valid loss:0.8874 valid f1:0.1518
Epoch:3 / 40, train loss:0.4129 train f1:0.2651, valid loss:0.8092 valid f1:0.3788
Epoch:4 / 40, train loss:0.3421 train f1:0.4463, valid loss:0.6782 valid f1:0.5184
Epoch:5 / 40, train loss:0.2919 train f1:0.5603, valid loss:0.5135 valid f1:0.6041
Epoch:6 / 40, train loss:0.2566 train f1:0.6342, valid loss:0.4853 valid f1:0.6382
Epoch:7 / 40, train loss:0.2315 train f1:0.6801, valid loss:0.4338 valid f1:0.6796
Epoch:8 / 40, train loss:0.2086 train f1:0.7119, valid loss:0.3959 valid f1:0.6945
Epoch:9 / 40, train loss:0.1907 train f1:0.7413, valid loss:0.3563 valid f1:0.7316
Epoch:10 / 40, train loss:0.1748 train f1:0.7681, valid loss:0.3440 valid f1:0.7205
Epoch:11 / 40, train loss:0.1610 train f1:0.7883, valid loss:0.3082 valid f1:0.7485
Epoch:12 / 40, train loss:0.1480 train f1:0.8048, valid loss:0.3221 valid f1:0.7592
Epoch:13 / 40, train loss:0.1361 train f1:0.8246, valid loss:0.2428 valid f1:0.7642
Epoch:14 / 40, train loss:0.1249 train f1:0.8395, valid loss:0.2354 valid f1:0.7580
Epoch:15 / 40, train loss:0.1149 train f1:0.8553, valid loss:0.2244 valid f1:0.7655
Epoch:16 / 40, train loss:0.1059 train f1:0.8679, valid loss:0.1923 valid f1:0.7702
Epoch:17 / 40, train loss:0.0949 train f1:0.8812, valid loss:0.2187 valid f1:0.7773
Epoch:18 / 40, train loss:0.0884 train f1:0.8930, valid loss:0.1699 valid f1:0.7858
Epoch:19 / 40, train loss:0.0787 train f1:0.9066, valid loss:0.1741 valid f1:0.7783
Epoch:20 / 40, train loss:0.0724 train f1:0.9145, valid loss:0.1458 valid f1:0.7902
Epoch:21 / 40, train loss:0.0674 train f1:0.9209, valid loss:0.1417 valid f1:0.7780
Epoch:22 / 40, train loss:0.0630 train f1:0.9289, valid loss:0.1141 valid f1:0.7927
Epoch:23 / 40, train loss:0.0580 train f1:0.9320, valid loss:0.1295 valid f1:0.7958
Epoch:24 / 40, train loss:0.0535 train f1:0.9392, valid loss:0.1321 valid f1:0.7924
Epoch:25 / 40, train loss:0.0490 train f1:0.9452, valid loss:0.0996 valid f1:0.7954
Epoch:26 / 40, train loss:0.0450 train f1:0.9501, valid loss:0.0774 valid f1:0.7967
Epoch:27 / 40, train loss:0.0417 train f1:0.9549, valid loss:0.0924 valid f1:0.8037

```

Epoch:28 / 40, train loss:0.0379 train f1:0.9601, valid loss:0.0903 valid f1:0.7953
 Epoch:29 / 40, train loss:0.0367 train f1:0.9612, valid loss:0.0712 valid f1:0.7898
 Epoch:30 / 40, train loss:0.0342 train f1:0.9641, valid loss:0.0571 valid f1:0.8014
 Epoch:31 / 40, train loss:0.0320 train f1:0.9673, valid loss:0.0751 valid f1:0.8034
 Epoch:32 / 40, train loss:0.0298 train f1:0.9705, valid loss:0.0548 valid f1:0.7976
 Epoch:33 / 40, train loss:0.0281 train f1:0.9718, valid loss:0.0609 valid f1:0.8036
 Epoch:34 / 40, train loss:0.0270 train f1:0.9725, valid loss:0.0461 valid f1:0.8048
 Epoch:35 / 40, train loss:0.0259 train f1:0.9739, valid loss:0.0540 valid f1:0.7946
 Epoch:36 / 40, train loss:0.0235 train f1:0.9768, valid loss:0.0429 valid f1:0.8055
 Epoch:37 / 40, train loss:0.0230 train f1:0.9777, valid loss:0.0411 valid f1:0.7991
 Epoch:38 / 40, train loss:0.0223 train f1:0.9772, valid loss:0.0360 valid f1:0.7951
 Epoch:39 / 40, train loss:0.0217 train f1:0.9785, valid loss:0.0426 valid f1:0.7927
 Epoch:40 / 40, train loss:0.0194 train f1:0.9813, valid loss:0.0341 valid f1:0.8020

Classification Report for CNN with accuracy of 81%

In [42]:

```
#ModelCNN
from sklearn.metrics import classification_report

print(classification_report(Y_test, preds))
```

	precision	recall	f1-score	support
0	0.95	0.70	0.80	191
1	0.97	0.78	0.87	171
2	0.94	0.83	0.88	194
3	0.97	0.85	0.90	134
4	0.93	0.88	0.91	104
5	0.92	0.88	0.90	264
6	0.92	0.78	0.85	187
7	0.91	0.83	0.86	219
8	0.87	0.64	0.74	112
9	0.85	0.54	0.66	177
10	0.91	0.69	0.78	102
11	0.92	0.79	0.85	147
12	0.83	0.65	0.73	344
13	0.94	0.86	0.89	229
14	0.98	0.94	0.96	223
15	0.96	0.85	0.90	109
16	0.97	0.78	0.87	218
17	0.86	0.73	0.79	101
18	0.95	0.85	0.89	165
19	0.84	0.59	0.69	73
20	0.86	0.64	0.73	261
21	0.94	0.86	0.90	149
22	0.80	0.50	0.62	175
23	0.83	0.75	0.79	259
24	0.83	0.74	0.78	184
25	0.88	0.82	0.85	178
26	0.93	0.70	0.80	101
27	0.93	0.83	0.88	203
28	0.96	0.87	0.91	138
29	0.86	0.65	0.74	319
30	0.93	0.68	0.79	152
31	0.85	0.60	0.70	169
32	0.97	0.75	0.85	89
33	0.93	0.91	0.92	199
34	0.92	0.80	0.86	200
35	0.85	0.24	0.37	143
36	0.95	0.91	0.93	111
37	0.88	0.60	0.71	112
38	0.85	0.73	0.79	304
39	0.83	0.72	0.77	179
40	0.83	0.60	0.70	134

41	0.94	0.50	0.66	121
42	0.81	0.80	0.80	172
43	0.94	0.94	0.94	204
44	0.88	0.73	0.80	192
45	0.95	0.76	0.85	123
46	0.85	0.80	0.82	308
47	0.90	0.88	0.89	152
48	0.97	0.79	0.87	157
49	0.90	0.86	0.88	144
micro avg	0.90	0.75	0.82	8796
macro avg	0.90	0.75	0.81	8796
weighted avg	0.90	0.75	0.81	8796
samples avg	0.86	0.74	0.78	8796

```
/usr/local/lib/python3.8/dist-packages/sklearn/metrics/_classification.py:1318: UndefinedMetricWarning: Precision and F-score are ill-defined and being set to 0.0 in samples with no predicted labels. Use `zero_division` parameter to control this behavior.  
  _warn_prf(average, modifier, msg_start, len(result))
```

In []: